

Assignment 1 Report | Characterising an AMOC timeseries

- Author: Hengxi Yang
- Date: Sep 2026
- Course : 63-731 Data Analysis in Physical Oceanography

1. Introduction

1.1 Aim

- Load the DSO transport timeseries and fill its NaN positions. Plot a Frequency distribution figure to illustrate the volume distribution.
- Compute and plot a power spectrum of DSO transport using Welch's overlapped-segment averaging. Furthermore, verify the variance budget with Parseval and apply a low-pass filter on the original timeseries, plot a timeseries and a power spectrum to compare the original one and the filtered one.
- Apply two different windows and see which one has a better ability to suppress the sidelobes. Add a chi-squared confidence band to the spectrum with EDF.

1.2 Data and Series

- **Dataset:** Denmark Strait Overflow, DSO, hourly
 - **Series:** Overflow time-series through Denmark Strait
 - **Timespan:** 1996-05-01 to 2021-08-07
 - **Record length:** 221,514
-

2. Results and Discussion

Part A Characterise the Series

The series I chose is Denmark Strait Overflow (DSO). I tried Calafat 2025 before, which includes a meridional heat transport timeseries with a sample size of only 66, it was too small for the following processing. I also tried OSNAP west, which also has a small sample size. Admittedly, DSO is the lower limb of the AMOC rather than direct overturning transport, but it is a core component of the AMOC. Furthermore, options are limited, and the DSO transport timeseries has a long timespan with 221,514 data points, making it an optimal choice.

Here is some info on trans_DSO:


Parameter	Value
Record Length	221,514 time steps

Parameter	Value
Timespan	1996-05-01 to 2021-08-07
Sampling Interval	~ 1 hour
Missing Values	18,456

Because there're a lot pf NaNs, I can't just drop them. So I used **linear interpolation** to fill the NaN positions. I didn't know if there's any pattern or period in Trans_DSO, so linear interpolation is the safest way. For the remaining NaNs at both ends, I dropped them directly. Here is some info on the new Tran_DSO.

Parameter	Value
Record Length	220,956 time steps
Timespan	1996-05-24 to 2021-08-07
Sampling Interval	~ 1 hour
Missing Values	0

The timespan shortened, but it was worth it.




Parameter	Value (Sv)
Mean	-3.125
Standard Deviation	1.457
Minimum	-9.661
Maximum	3.084

Parameter	Value (Sv)
Range	12.745

Bin width = $\frac{12.745 \text{ Sv}}{100} = \mathbf{0.12745 \text{ Sv}}$

According to the distribution figure, the southward transport concentrates around the mean, which indicates that the DSO might have some specific flow characteristics. The standard deviation $\sigma = 1.46 Sv$ indicates that the transport timeseries has significant fluctuations.

Part B The Spectrum




Parameter	Value
segment length	100-day seasonal window
noverlap	50% (default)
sample frequency	24
Nyquist frequency	12 cycles/day

The PSD figure shows that most of power concentrates around the **low frequency area** . There's a steep decrease from 10⁻¹ to 10⁰ cycles/day, which means the power of DSO distributes in seasonal and longer time scales.

Additionally, there's a significant peak around 2 cycles/day, which can be seen clearly in the semi-log figure:

Welch's Power Spectral Density of DSO



Denmark Strait is situated between Greenland and Iceland, where semidiurnal tides (such as the M2) dominate. The period of M2 is about half of a day, so most of the energy is focused around 2 cycles/day.



The red noise spans from 10^{-2} to $\{10^0\}$ cycles/day, the low frequency range reveals long-term signal, while in the later area, the higher the frequency, the lower the PSD.

The white noise starts from ~ 4 cycles/day, where the PSD keeps nearly constant across frequencies, indicating that the real oceanographic signals have decayed.

The Parseval's Theorem is:


$$\text{Variance}_{\text{time}} \approx \sum \text{PSD} \times \Delta f$$

The verification results:

Parameter	Value
Time-domain Variance	2.122 Sv ²
Frequency-domain Integral (PSD)	1.931 Sv ²
Difference	0.191 Sv ²
Relative Error	8.992%
Tolerance	10%
Test Result	Passed

The relative error is 8.992%, clearly passes the 10% threshold.


After applying 30-day Tukey low-pass window, the amplitude of Trans_SDO narrows considerably.



Due to the data gaps around 2000 and 2007, the Tukey Filter cannot smooth the signals properly.


The 30-day low-pass filter removed a lot rapid and strong fluctuations.

PSD Comparison of Original and Low-Pass Filtered DSO




From 10^{-1} to 10^1 cycles/day, the filtered PSD shows periodic ripples. Overall, the filtered PSD has a bigger slope relative to the original one. As high-frequency signal (white noise) are removed, the power of the filtered one is significantly lower than the original one starting from around 2 cycles/day (high-frequency area), which can be seen clearly in the semilog figure:

PSD Comparison of Original and Low-Pass Filtered DSO




Part C Filter Design and Spectral Confidence Interval

I applied two different windows on the DSO transport series, one is Tukey, and another one is a pure box without any edge smoothing, both are 30-day low-pass filters. I computed the frequency responses after normalization, here is the result.




Without any tapering, the boxcar window exhibits periodic sidelobes, leaving more high-frequency energy at the same time. In contrast, the Tukey window shows a better ability to attenuate the high-frequency signals.

The following figure illustrates the Welch's PSD with Chi-squared Confidence interval.



If we zoom in and focus on the salient peak:



The peak lies within the 95% confidence band, but it rises sharply above the background noise, making it a highly significant tidal signal.